

Programming Assignment: Fast File Duplication

Operating Systems@NYCU

Prof. Li-Pin Chang

Objectives

- Understand the overheads involved in file operations
- Exploit kernel primitives to avoid (or to reduce) these overheads

The Baseline Method

- open → (read, write)*N

```
while(1){
    readn = read(fd_in, buf, BUF_SIZE);
    if(readn <= 0){
        break;
    }
    while(readn){
        size_t writen = write(fd_out, buf, readn);
        if(writen == -1){
            perror("write() failed");
            return -1;
        }
        readn -= writen;
    }
}
```

- Can you do better (faster) than this?

Overheads of File Rding(Wrting)

- Switching from the user mode to the kernel mode
 - Copying data from the storage to kernel pages
 - Copying data from the kernel pages to user buffer
 - Switching from the kernel mode to the user mode
-
- In this assignment, we duplicate a huge file (1GB) on *tmpfs*, a RAM-based file system, to focus on the OS-level overhead rather than the storage latency

Test Procedure

- Pre-conditioning
 - mount tmpfs: ``mount -t tmpfs -o size=3G tmpfs <mountpoint>``
 - clear swap: ``swapoff -a && swapon -a``
 - drop cache: ``echo 3 > /proc/sys/vm/drop_caches``
- Run your program
- Cleaning-up
 - diff: “diff source destination”
- Report the time spent on the 2nd step
 - The diff result must report **no difference**
 - ****No need** to do `fsync()` on the destination file

Reference time

- TA's time to duplicate a 1GB file on tmpfs

Method	Baseline	Method 1	Method 2	Method 3	Method 4
Time (s)	0.755	0.350	0.622	0.287	0.255
Time	100%	46%	82%	38%	34%

* The time may vary depending on your hardware configuration.

- TA's test platform
 - Ubuntu 20.04.3 LTS
 - Intel Core i5-10400 with 32GB ram

Requirement

- Implement your own method of file duplication
- Your method must be *faster* than Baseline
- **Do not use `link()` to “duplicate” the file**
- Turn in a 2-page report to E3 (pdf)
 - Compare the performance of yours vs. baseline
 - Describing all the implementation details
 - Explain why your method improves the performance
- Turn in your program to E3 (.zip, .c, or .cxx)
 - TA will compile your program and verify the performance of your method

Template of Your 2-Page Report

1. Section 1: Analysis of Baseline
 - This section describes the performance bottleneck of the baseline method
2. Section 2: Proposed method
 - This section describes your proposed method
3. Section 3: Performance evaluation
 - Present the measured performance results of Baseline and your method
 - Discuss why and how your method will perform better than. Baseline
4. References

How to Get a High Mark

- Grading Policy
 - 50% quality of your report
 - 50% time efficiency of your method
- (Program) Duplicate the file correctly and considerably improve upon Baseline
- (Report) Provide insightful discussion, not just show the raw numbers

Source Code of Baseline

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>

#define BUF_SIZE 1024

int main(int argc, char *argv[]){
    if(argc != 3){
        printf("Usage: %s <source> <destination>\n", argv[0]);
        return -1;
    }

    int fd_in = open(argv[1], O_RDONLY);
    if(fd_in == -1){
        perror("readwrite: open");
        return -1;
    }

    int fd_out = creat(argv[2], 0644);
    if(fd_out == -1){
        perror("readwrite: creat");
        return -1;
    }
}
```

```
char buf[BUF_SIZE];
size_t readn;
while(1){
    readn = read(fd_in, buf, BUF_SIZE);
    if(readn <= 0){
        break;
    }
    while(readn){
        size_t writen = write(fd_out, buf, readn);
        if(writen == -1){
            perror("readwrite: write");
            return -1;
        }
        readn -= writen;
    }
}

if(readn == -1){
    perror("readwrite: read");
    return -1;
}

close(fd_in);
close(fd_out);
return 0;
}
```

Deadline

23:55, June 17, 2022

Turn-in all your materials to E3